

HW2

October 5, 2025

1 HW 2

This assignment covers several aspects of Linear Regression. **DO NOT ERASE MARKDOWN CELLS AND INSTRUCTIONS IN YOUR HW submission**

- Q - QUESTION
- A - Where to input your answer

1.1 Instructions

Keep the following in mind for all notebooks you develop: * Structure your notebook. * Use headings with meaningful levels in Markdown cells, and explain the questions each piece of code is to answer or the reason it is there. * Make sure your notebook can always be rerun from top to bottom (Kernel Tab -> Restart and Run All) * Start working on this assignment as soon as possible. If you are a beginner in Python, this might take a long time. One of the objectives of this assignment is to help you learn Python and the sci-kit package. * Follow [README.md](#) for homework submission instructions * In this notebook, we assume './data/' location of all data files to be read and written

1.2 Related sklearn material and online tutorials

[sklearn User Guide](#)

1.2.1 sklearn data pre-processing

- [train_test_split](#)
- [common_pitfalls](#)
- [train test split tutorial](#)

1.2.2 sklearn multiple linear regression

- [tutorial](#)
- [API documentation](#)
- [Linear Regression](#)
- [multiple linear regression tutorial](#)

1.2.3 sklearn polynomial regression

- [generate polynomial features](#)
- [polynomial regression tutorial](#)

1.2.4 correlation

- [correlation](#)

2 Linear Regression

In jupyter notebook environment, commands starting with the symbol % are magic commands or magic functions. %%timeit is one of such function. It basically gives you the speed of execution of certain statement or blocks of codes.

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

Q1 Read the car_data.csv data (we assume ../data/ location of all data files to be read and written) from data folder using pandas. Replace the ??? in the code cell below to accomplish this task.

A1 Replace ??? with code in the code cell below

```
[7]: # Replace ??? with code in the code cell below

df = pd.read_csv('../data/Car_data.csv')
```

```
[12]: # View head of the data to confirm the correctness of your answer
df.head()
```

```
[12]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	\
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

	carbody	drivewheel	engine	location	wheelbase	...	enginesize	\
0	convertible	rwd	front	88.6	...	130		
1	convertible	rwd	front	88.6	...	130		
2	hatchback	rwd	front	94.5	...	152		
3	sedan	fwd	front	99.8	...	109		
4	sedan	4wd	front	99.4	...	136		

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

[5 rows x 26 columns]

2.1 Data cleaning and manipulation

Q2 Here, you will practice the usage of common data cleaning and manipulation functions in 3 steps. 1. Use `isnull()` to figure out the number of NaN values per column 2. Remove the column with majority NaN values (if any), else make comment with `#` in front of the cell 3. Check if there are still NaN values in the dataframe using `isna()` method

A2 Replace ??? with code in the code cell below

```
[13]: # Is there any missing data here on this dataset :
df.isnull()
```

```
[13]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
..	
200	False	False	False	False	False	False	False	
201	False	False	False	False	False	False	False	
202	False	False	False	False	False	False	False	
203	False	False	False	False	False	False	False	
204	False	False	False	False	False	False	False	

	drivewheel	enginelocation	wheelbase	...	enginesize	fuelsystem	\
0	False	False	False	...	False	False	
1	False	False	False	...	False	False	
2	False	False	False	...	False	False	
3	False	False	False	...	False	False	
4	False	False	False	...	False	False	
..	
200	False	False	False	...	False	False	
201	False	False	False	...	False	False	
202	False	False	False	...	False	False	
203	False	False	False	...	False	False	
204	False	False	False	...	False	False	

	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
--	-----------	--------	------------------	------------	---------	---------	---

0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
..
200	False	False	False	False	False	False
201	False	False	False	False	False	False
202	False	False	False	False	False	False
203	False	False	False	False	False	False
204	False	False	False	False	False	False

	highwaympg	price
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
..
200	False	False
201	False	False
202	False	False
203	False	False
204	False	False

[205 rows x 26 columns]

```
[14]: #Remove the column with majority NaN values (if any)
df.dropna(axis=1, thresh=len(df)//2+1)
```

```
[14]:      car_ID  symboling      CarName fueltype aspiration \
0         1         3    alfa-romero giulia      gas      std
1         2         3    alfa-romero stelvio      gas      std
2         3         1  alfa-romero Quadrifoglio      gas      std
3         4         2        audi 100 ls      gas      std
4         5         2        audi 100ls      gas      std
..      ...      ...      ...      ...      ...
200      201        -1    volvo 145e (sw)      gas      std
201      202        -1    volvo 144ea      gas    turbo
202      203        -1    volvo 244dl      gas      std
203      204        -1    volvo 246    diesel    turbo
204      205        -1    volvo 264gl      gas    turbo

      doornumber      carbody drivewheel enginelocation  wheelbase  ... \
0          two  convertible      rwd      front      88.6  ...
1          two  convertible      rwd      front      88.6  ...
2          two   hatchback      rwd      front      94.5  ...
```

3	four	sedan	fwd	front	99.8	...
4	four	sedan	4wd	front	99.4	...
..
200	four	sedan	rwd	front	109.1	...
201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...

	engine	size	fuel	system	bore	ratio	stroke	compression	ratio	horsepower	\
0		130	mpfi		3.47	2.68		9.0		111	
1		130	mpfi		3.47	2.68		9.0		111	
2		152	mpfi		2.68	3.47		9.0		154	
3		109	mpfi		3.19	3.40		10.0		102	
4		136	mpfi		3.19	3.40		8.0		115	
..		
200		141	mpfi		3.78	3.15		9.5		114	
201		141	mpfi		3.78	3.15		8.7		160	
202		173	mpfi		3.58	2.87		8.8		134	
203		145	idi		3.01	3.40		23.0		106	
204		141	mpfi		3.78	3.15		9.5		114	

	peak	rpm	city	mpg	highway	mpg	price
0		5000	21		27		13495.0
1		5000	21		27		16500.0
2		5000	19		26		16500.0
3		5500	24		30		13950.0
4		5500	18		22		17450.0
..	
200		5400	23		28		16845.0
201		5300	19		25		19045.0
202		5500	18		23		21485.0
203		4800	26		27		22470.0
204		5400	19		25		22625.0

[205 rows x 26 columns]

```
[15]: # lets get some statistical information :
df.describe()
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	

75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000

	curbweight	enginesize	boreratio	stroke	compressionratio	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	2555.565854	126.907317	3.329756	3.255415	10.142537	
std	520.680204	41.642693	0.270844	0.313597	3.972040	
min	1488.000000	61.000000	2.540000	2.070000	7.000000	
25%	2145.000000	97.000000	3.150000	3.110000	8.600000	
50%	2414.000000	120.000000	3.310000	3.290000	9.000000	
75%	2935.000000	141.000000	3.580000	3.410000	9.400000	
max	4066.000000	326.000000	3.940000	4.170000	23.000000	

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

Q3: In this task, out of all categorical columns, we focus only on the **fueltype** column processing in 2 steps. 1. Use label encoder from sklearn and convert the **fueltype** categorical values to numerical values.

2. Create a new dataframe that contains only the numerical columns.

A3 Replace ??? with code in the code cell below.

```
[17]: # Label Encoding for 2-class columns:
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['fueltype'] = le.fit_transform(df['fueltype'])
```

```
[18]: # Create new dataframe with selected columns. Careful, don't rename ex:
      ↪ df_numeric
df=df.select_dtypes(include='number')
```

```
[19]: df.head()
```

```
[19]:   car_ID  symboling  fueltype  wheelbase  carlength  carwidth  carheight  \
0        1          3         1         88.6         168.8         64.1         48.8
1        2          3         1         88.6         168.8         64.1         48.8
2        3          1         1         94.5         171.2         65.5         52.4
3        4          2         1         99.8         176.6         66.2         54.3
4        5          2         1         99.4         176.6         66.4         54.3
```

	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	\
0	2548	130	3.47	2.68	9.0	111	
1	2548	130	3.47	2.68	9.0	111	
2	2823	152	2.68	3.47	9.0	154	
3	2337	109	3.19	3.40	10.0	102	
4	2824	136	3.19	3.40	8.0	115	

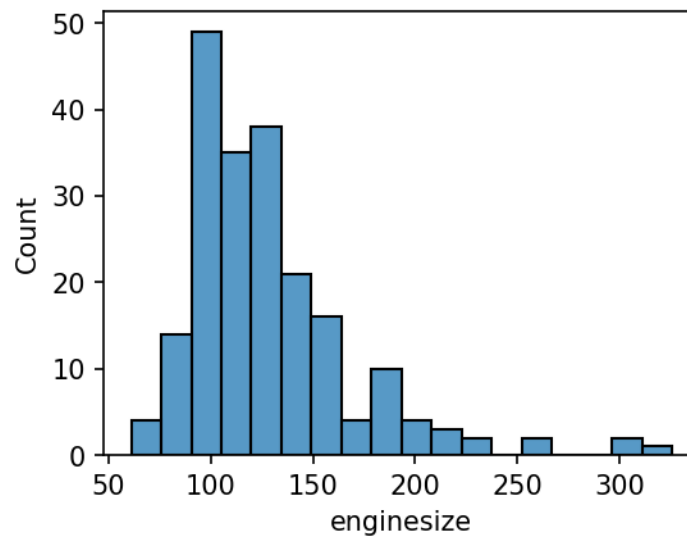
	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0

Q4: Use seaborn histplot to plot a distribution graph for the engine sizes

A4 Replace ??? with code in the code cell below

```
[21]: plt.figure(figsize=(4,3),dpi=150)
      sns.histplot(df['enginesize'])
```

```
[21]: <Axes: xlabel='enginesize', ylabel='Count'>
```

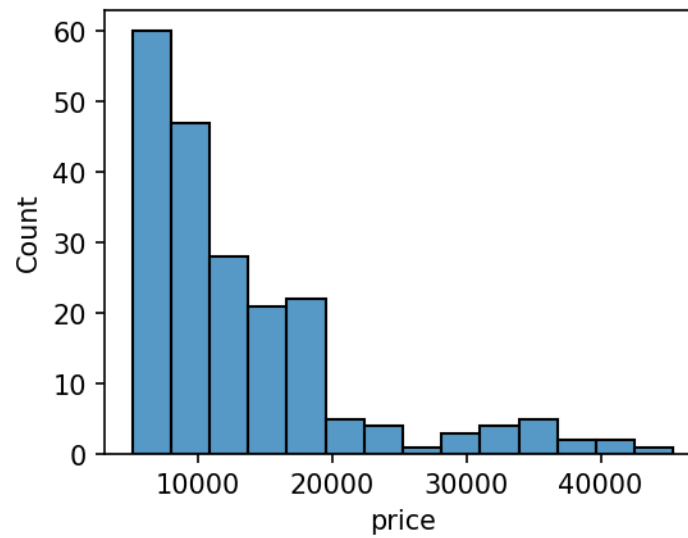


Q5: Use seaborn histplot to plot a distribution graph for the car prices

A5 Replace ??? with code in the code cell below

```
[24]: plt.figure(figsize=(4,3),dpi=150)
      sns.histplot(df['price'])
```

```
[24]: <Axes: xlabel='price', ylabel='Count'>
```

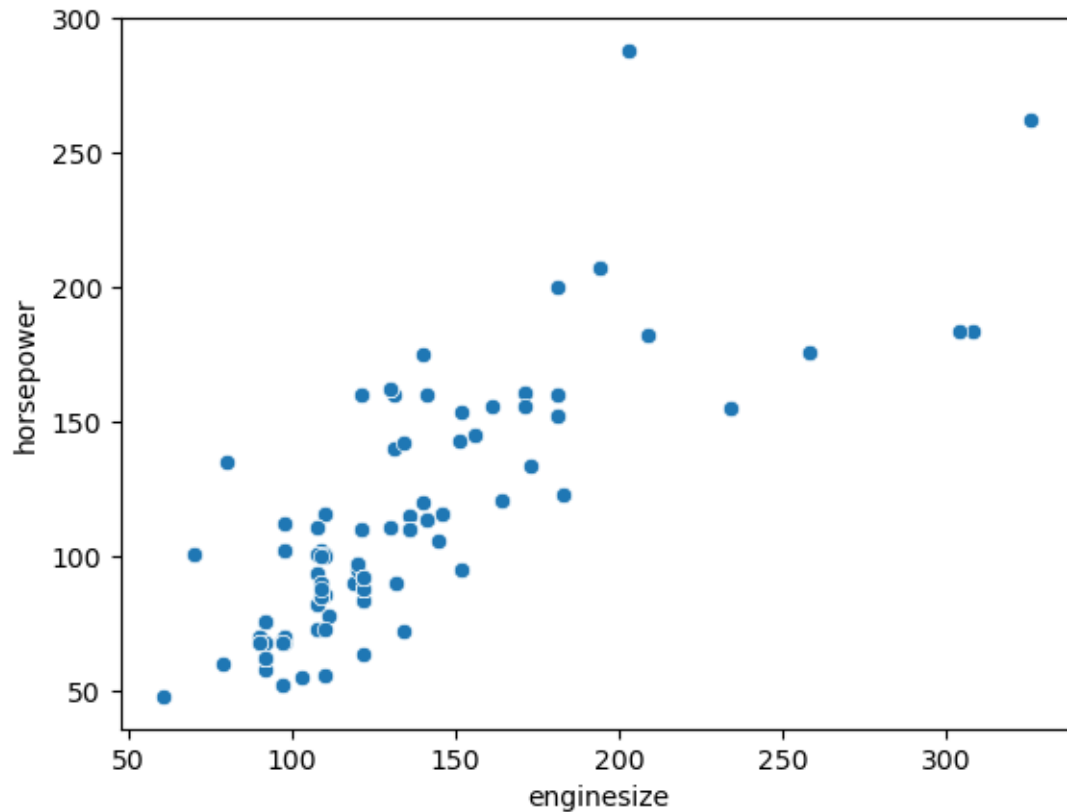


Q6: Use seaborn scatterplot to present the relation between enginesize and the horsepower of a car

A6 Replace ??? with code in the code cell below

```
[25]: sns.scatterplot(x='enginesize', y='horsepower', data=df)
```

```
[25]: <Axes: xlabel='enginesize', ylabel='horsepower'>
```

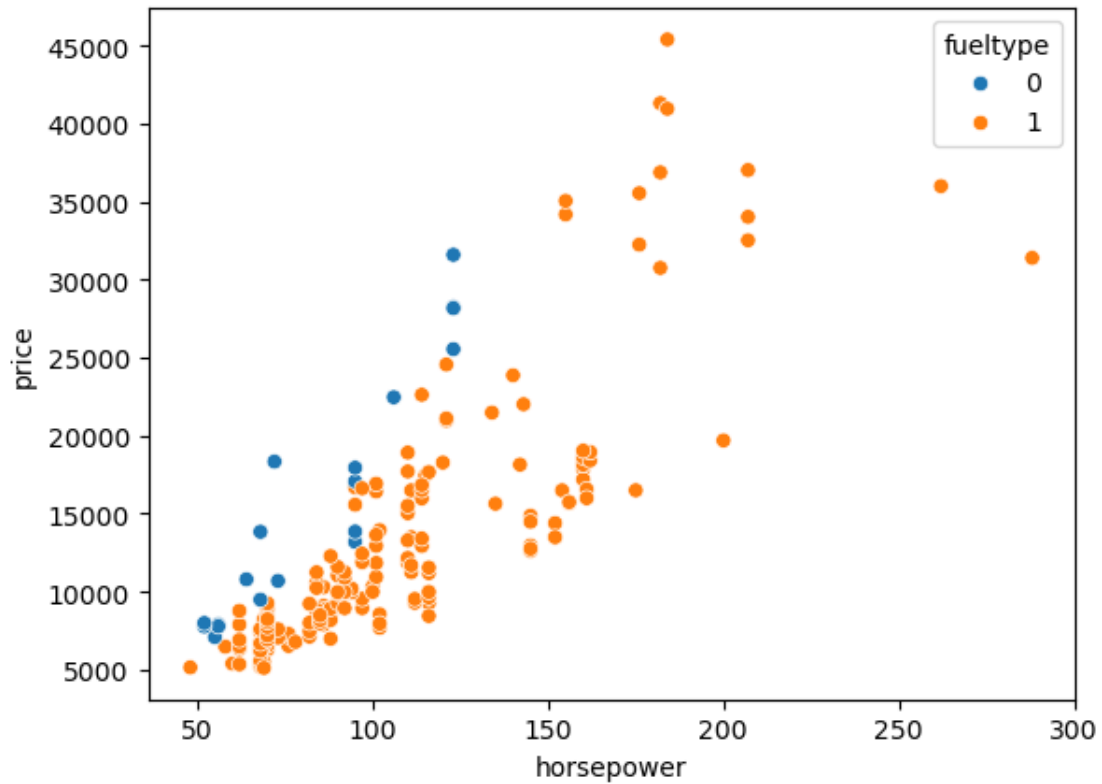
Q7: There is a correlation between the car price and the horsepower of a car. If horsepower of a car increase, the price of the car also increases most of the time, and in this question you will use the seaborn scatterplot to present the relation between price and horsepower.

Next, use `hue` parameter of scatterplot function to illustrate datapoints that relate to specific fueltype category.

A7 Replace ??? with code in the code cell below

```
[26]: sns.scatterplot(x='horsepower', y='price', hue='fueltype', data=df)
```

```
[26]: <Axes: xlabel='horsepower', ylabel='price'>
```

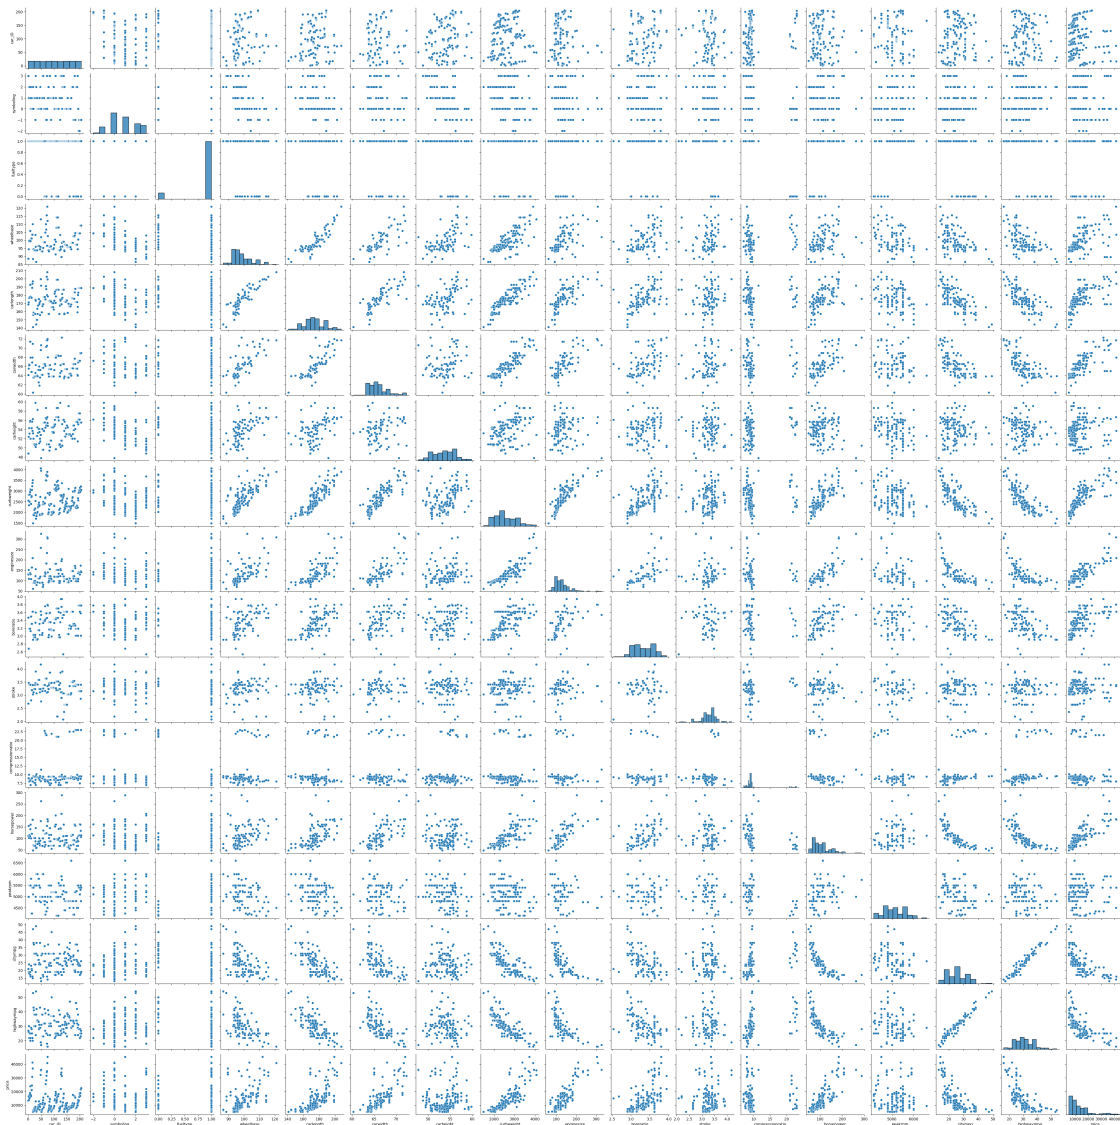


Q8: Use pairplot from sns to plot the data frame `df` and justify your feature selection.

A8: replace ??? with code in the code cell below.

```
[27]: # 2. Use pairplot from sns to plot our data frame df
sns.pairplot(df)
```

```
[27]: <seaborn.axisgrid.PairGrid at 0x32c23d400>
```



Q9 Data Visualization:

1. Use heatmap chart from seaborn library to findout the correlation between the columns in our dataset.
2. Update data frame 'df' to contain 5 columns from existing 'df' with the highest correlation to column "price". Also include price column in the updated data frame.

A9 Replace ??? with code in the code cell below

```
[28]: corr_matrix = df.corr()
plt.figure(figsize=(14,14))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm")
```

[28]: <Axes: >


```
[30]: y =df['price']
X =df.drop('price', axis=1)
X
```

```
[30]:      enginesize  curbweight  horsepower  carwidth  highwaympg
0           130         2548         111      64.1         27
1           130         2548         111      64.1         27
2           152         2823         154      65.5         26
3           109         2337         102      66.2         30
4           136         2824         115      66.4         22
..          ...          ...          ...      ...          ...
200          141         2952         114      68.9         28
201          141         3049         160      68.8         25
202          173         3012         134      68.9         23
203          145         3217         106      68.9         27
204          141         3062         114      68.9         25
```

[205 rows x 5 columns]

Q11 Use `train_test_split` to split the data set as train:test=(75%:25%) ratio.

A11 Replace ??? with code in the code cell below

```
[32]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳random_state=42)
# View the shape of your data set
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[32]: ((153, 5), (52, 5), (153,), (52,))
```

2.3 Regression Task

2.3.1 Multiple Linear Regression

Q12 Fit multiple linear regression model on training data using all predictors, see (i) [Linear Regression Example](#); (ii) [scikit-learn linear model](#)

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

A12: Replace ??? with code in the code cell below

```
[34]: from sklearn.linear_model import LinearRegression
linear_model = LinearRegression()
linear_model.fit(X_train, y_train)
```

```
[34]: LinearRegression()
```

Q13: Model Scoring 1. Calculate the test MSE 2. Print the score from the model using test data

A13 Replace ??? with code in the code cell below

```
[35]: # Calculate the score on train and test sets
# Your code goes below
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
y_pred=linear_model.predict(X_test)
mse = mean_squared_error(y_test, y_pred) # Calculate the test MSE
print("Test mean squared error (MSE): {:.2f}".format(mse))

print(y_pred)
```

```
Test mean squared error (MSE): 14041064.92
[26090.89300501 19041.96831161 11282.01510357 13768.7747931
 23417.84670716  6488.8582937   6731.59600438  7375.70031489
 10692.93716198  6236.67774343 15565.73401052  7191.84067183
 15673.55244404 12304.60929575 38563.28129637  5491.50104535
 -1681.84856428 18842.11863788 11376.68555448 10481.56097807
 11879.06955665 21827.73930583  6465.95776328  3852.90908061
  5743.67538258 26905.50033411 15236.15410078 16483.23560903
  6499.16353239 16343.54237348 23148.76547475  5718.48479912
  6284.57214268 21599.31136587  8998.34450504 23125.86494434
 11809.84910882  8721.27893361  5230.2518289  18948.60610432
  9792.87020249 11738.73172882 14915.9927592  5874.6312074
  6423.59178201  9956.28155646  5718.48479912  8047.19820906
 16998.9051383  18816.92805442  5190.17590067 21965.36164919]
```

2.3.2 Polinomial Regression

Q14: Polynomial extension of the feature set captures the non-linear dependencies in the data 1. Create a polinomial feature transformer with degree **TWO** using sklearn library [PolynomialFeatures](#) 2. Transform the training dataset using the polinomial feature transformer

A14 Replace ??? with code in the code cell below

```
[36]: from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=2)
poly_features = poly.fit_transform(X_train)
```

Q15: Train the new model 1. Create a LinearRegression model using sklearn 2. Train the model using the transformed Train data(X_train)/ or Polinomial train data 3. Print the score for the Polinomial Regression for the Train data.

See (i) [Linear Regression Example](#); (ii) Use the transformed X_train features inside the score() function for the correct model scores.

A15 Replace ??? with code in the code cell below

```
[38]: poly_reg_model = LinearRegression()
poly_reg_model.fit(poly_features, y_train)
```

```
poly_reg_model.score(poly_features, y_train)
```

```
[38]: 0.8928327396853508
```